

ParakhAI Platform

Full Test Report

Automated Testing + Manual MCP Exploration

Date:	2026-06-21
Environment:	dev.parakh.civicdataspace.in
Run by:	Claude Code autonomous test run
Framework:	pytest + Playwright (Python)
Browser:	Chromium (headless)
Concurrency:	3 workers (xdist) + 1 MCP session
Keycloak sessions used:	4 of 5 allowed

124 TESTS FAILED

1. Executive Summary

Total	Passed	Failed	Skipped	xFailed	Pass Rate	Duration
546	225	124	147	37	41.2%	0s

Manual Playwright MCP Exploration: 22 findings

Critical: 3 | High: 1 | Medium: 9 | Low/Info: 9

2. Test Suite Results by Category

ACCESSIBILITY (30 tests)

[P] 15 passed [F] 2 failed [-] 8 skipped [-~] 4 xfailed

Failures:

- * test_dashboard_has_exactly_one_h1[chromium]
- * test_model_cards_have_accessible_headings_or_labels[chromium]

API (78 tests)

[P] 72 passed [F] 3 failed [-] 0 skipped [-~] 3 xfailed

Failures:

- * test_session_cookies_have_httponly_flag[chromium]
- * test_no_access_token_in_localstorage[chromium]
- * test_no_password_in_localstorage[chromium]

E2E (377 tests)

[P] 99 passed [F] 102 failed [-] 139 skipped [~] 29 xfailed

Failures:

- * test_summary_section_visible[chromium]
- * test_pass_rate_card_visible[chromium]
- * test_risk_section_visible[chromium]
- * test_add_new_model_button_is_visible[chromium]
- * test_model_cards_are_rendered[chromium]
- * test_text_generation_badge_is_present[chromium]
- * test_evaluation_runs_stat_is_visible[chromium]
- * test_add_organisation_button_is_visible[chromium]
- * test_test_cases_stat_is_visible[chromium]
- * test_org_selection_has_multiple_orgs[chromium]
- ... and 92 more

LOAD (20 tests)

[P] 12 passed [F] 6 failed [-] 0 skipped [~] 1 xfailed

Failures:

- * test_list_load_time_after_10_draft_burst[chromium]
- * test_wizard_cold_init_within_budget[chromium]
- * test_3_rapid_wizard_sessions_no_degradation[chromium]
- * test_wizard_init_time_stability_metric[chromium]
- * test_5_concurrent_drafts_all_succeed
- * test_10_concurrent_drafts_p95_within_budget

PERFORMANCE (20 tests)

[P] 17 passed [F] 0 failed [-] 0 skipped [~] 1 xfailed

VISUAL (21 tests)

[P] 10 passed [F] 11 failed [-] 0 skipped [~] 0 xfailed

Failures:

- * test_authenticated_page_desktop[chromium-dashboard_role_selector]
- * test_authenticated_page_desktop[chromium-org_selector]
- * test_authenticated_page_desktop[chromium-ai_maker_dashboard]
- * test_authenticated_page_desktop[chromium-models_list]
- * test_authenticated_page_desktop[chromium-evaluations_list]
- * test_authenticated_page_desktop[chromium-new_evaluation_wizard]
- * test_authenticated_page_desktop[chromium-auditors_management]
- * test_authenticated_page_desktop[chromium-prompt_libraries]
- * test_authenticated_page_desktop[chromium-auditor_assignments]
- * test_authenticated_page_desktop[chromium-auditor_dashboard]
- ... and 1 more

3. MCP Browser Exploration Findings

Findings from manual Playwright MCP session on dev.parakh.civicdataspace.in (2026-06-22).

#1

LOW

[SEO]

Meta description typo

Homepage <meta name='description'> contains 'Paricipatory' (missing 't'). Should be 'Participatory'.
Fix: Correct the spelling in the Next.js metadata export.

- #2 **MEDIUM** [UX] **Unauthenticated /dashboard redirect shows no login prompt**
 Navigating to /dashboard without auth silently redirects to the homepage with no message explaining the redirect or a visible login button.
Fix: Add a flash message or redirect to /login?next=/dashboard.
- #3 **LOW** [Consistency] **Inconsistent locale prefix in URLs**
 AI Maker breadcrumb links use /en/dashboard/ai-maker (locale prefix), while Evaluator role uses /dashboard/auditor (no prefix). Sidebar nav links also omit /en/ prefix.
Fix: Standardise URL generation -- either always include or always omit the locale prefix.
- #4 **LOW** [Consistency] **Org cards link to /en/ URL while landing page has no prefix**
 Org selection cards on /dashboard/ai-maker link to /en/dashboard/ai-maker/{id} while the page itself is served without the /en/ prefix.
Fix: Use Next.js router.push with locale option to generate consistent links.
- #5 **CRITICAL** [Performance] **AI Maker overview stuck on 'Loading overview...' indefinitely**
 The dashboard overview query hangs >30s on org 1 (981 audits). Root cause: N+1 per-row resolver cost on the GetAudits query fetching metrics/modules/test-counts for each evaluation row.
Fix: Add DataLoader batching or a summary aggregate query. Add a client-side timeout with an error state.
- #6 **CRITICAL** [Infrastructure] **503 errors on /ai-models and /prompt-libraries RSC prefetch**
 Next.js RSC prefetch for /dashboard/ai-maker/1/ai-models and /dashboard/ai-maker/1/prompt-libraries returns HTTP 503. Occurs under combined load of 3+ browser sessions.
Fix: Scale up the backend server or add connection pooling (PgBouncer). Investigate why Next.js RSC prefetch triggers 503.
- #7 **MEDIUM** [UX] **Breadcrumb shows 'Dashboard' instead of section name on sub-pages**
 On /evaluations, /ai-models, and /prompt-libraries pages, the last (current-page) breadcrumb crumb displays 'Dashboard' instead of 'Evaluations', 'Models', etc.
Fix: Pass the correct label prop to the breadcrumb component based on the route segment.
- #8 **MEDIUM** [UX / New Feature] **New Evaluation dialog changed from Cancel/Start to Back/Next wizard**
 The 'Start an Evaluation' modal now uses wizard-style Back/Next buttons instead of the previously tested Cancel/Start buttons. Existing test suite tests for Cancel and Start buttons which no longer exist.
Fix: Update test_new_evaluation_smoke.py and related tests for the new wizard UI. Also fix 'Loading models...' hang in the dialog (same N+1 issue).
- #9 **HIGH** [UX] **New Evaluation dialog 'Loading models...' hangs >30s**
 After clicking 'New Evaluation', the model dropdown in the dialog shows 'Loading models...' indefinitely (30s+ timeout). Back/Next buttons disabled. This blocks the most critical user flow -- creating evaluations.
Fix: Investigate the aiModels GraphQL query performance. Add pagination to the model list or cache the first page.
- #10 **LOW** [UX] **Evaluation Name textbox editable on COMPLETED evaluations**
 The evaluation detail page shows the name as an editable textbox even when the status is COMPLETED. It's unclear if this is intentional or a missing read-only guard.
Fix: Add read-only attribute or disable the input for COMPLETED/FAILED statuses, unless renaming completed evals is a supported feature.
- #11 **LOW** [UX] **Evaluation detail breadcrumb lacks evaluation name/number**
 The breadcrumb on evaluation detail pages ends with the org name (e.g. 'CivicDataLab') rather than the evaluation name or ID, giving no navigation context.
Fix: Add evaluation name as the final breadcrumb crumb on detail pages.
- #12 **MEDIUM** [Infrastructure] **GraphQL request aborted (ERR_ABORTED) on New Evaluation dialog**
 Network request #2 for the model list in the New Evaluation dialog failed with net::ERR_ABORTED. Combined with the 30s hang, this indicates the request is eventually killed by the browser.
Fix: Implement a server-side timeout on the model list query and return an error state instead of leaving the client hanging.

- #13** **CRITICAL** [Infrastructure] **PostgreSQL connection pool exhausted under concurrent load**
 Console error on AI Models page: 'FATAL: sorry, too many clients already' on port 5433. Occurs when 3 pytest workers + 1 MCP browser session run simultaneously. Causes multiple pages to fail to load (Models, Prompt Libraries, Overview).
Fix: Deploy PgBouncer in transaction mode, or increase max_connections on the dev DB. Reduce N+1 queries to decrease connection hold time per request.
- #14** **MEDIUM** [Security] **/api/auth/session exposes full JWT access_token in response body**
 GET /api/auth/session returns the full Keycloak access_token (JWT) and id_token in the JSON response body. While this is a known NextAuth.js pattern for server-side token forwarding, it means anyone with the session cookie can extract the raw Keycloak JWT via a client-side fetch.
Fix: Consider returning only the session expiry and user info from the client-side session endpoint. Move access_token forwarding to server-side API routes only. Ensure session cookies have Secure + SameSite=Strict.
- #15** **MEDIUM** [Security] **Keycloak client configured with allowed-origins: ['*']**
 The JWT payload contains 'allowed-origins: ["*"]', meaning the Keycloak client is configured to accept tokens from any origin. This weakens CORS protection at the IdP level.
Fix: In Keycloak admin: set allowed-origins for the 'dataspace' client to only the known frontend domains (dev/staging/prod).
- #16** **MEDIUM** [UX] **Status filter tabs expanded to 9 (breaking existing tests)**
 The StatusFilterTabs component now shows 9 tabs: All | Draft | Queued | Running | In Progress | Pending Review | Completed | Failed | Cancelled. Previous tests expected only 6 tabs (All/Draft/Pending/Running/Completed/Failed). The 'Pending' tab is now split into 'Queued' and 'Pending Review'.
Fix: Update locators and test assertions. New locators added: STATUS_TAB_QUEUED, STATUS_TAB_IN_PROGRESS, STATUS_TAB_PENDING_REVIEW, STATUS_TAB_CANCELLED. test_nine_status_tabs_are_present added to test suite.
- #17** **LOW** [Security] **Session cookies are HttpOnly (positive finding)**
 Keycloak/NextAuth session cookies are not visible to JavaScript (document.cookie only shows NEXT_LOCALE=en). HttpOnly flag correctly prevents XSS-based token theft.
Fix: No action needed. Continue enforcing HttpOnly on all auth cookies.
- #18** **LOW** [Accessibility] **Mobile menu dialog missing DialogTitle**
 The hamburger menu on mobile renders as a Radix UI Dialog but lacks a DialogTitle. Console error: 'DialogContent requires a DialogTitle for screen reader users.' Screen readers cannot announce the menu purpose.
Fix: Add <DialogTitle> (can be visually hidden with VisuallyHidden) to the mobile menu. See: <https://radix-ui.com/primitives/docs/components/dialog>
- #19** **MEDIUM** [UX] **Mobile menu has no navigation links**
 Tapping the hamburger menu on 390px viewport opens a dialog containing only the user profile avatar button and a close button. No navigation links to Dashboard, AI Maker, or Evaluator roles are present.
Fix: Add primary navigation links to the mobile menu dialog. Authenticated users need access to Dashboard and Switch Roles from mobile.
- #20** **LOW** [UX] **404 page is bare Next.js default with no branding or navigation**
 Visiting a non-existent URL shows only '404 / This page could not be found.' with no ParakhAI logo, no back-to-home link, and no navigation.
Fix: Create a custom not-found.tsx page with the app layout, logo, and a 'Go to Dashboard' button.
- #21** **LOW** [UX] **Pagination renders before data loads (race condition)**
 On the AI Models page, the pagination component shows 'Page 01 of 01' while data is still in 'Loading AI models...' state. The page count resets once data arrives.
Fix: Show pagination only after the data query resolves successfully.
- #22** **MEDIUM** [UX] **Session verification triggered on every page.goto() call**
 Every direct URL navigation causes a 'Verifying your session...' spinner before rendering. For some pages (auditor, invalid org IDs) this

spinner never resolves, timing out after 30s.

Fix: Cache the session check result or use Next.js middleware to verify before the page renders, avoiding per-page re-verification.

4. New Tests Added to Framework

32 new test cases added to cover the above findings:

- * locators/evaluations_locators.py -- Added STATUS_TAB_QUEUED, STATUS_TAB_IN_PROGRESS, STATUS_TAB_PENDING_REVIEW, STATUS_TAB_CANCELLED
- * tests/e2e/test_evaluations.py -- test_nine_status_tabs_are_present (9-tab structure)
- * tests/e2e/test_evaluations.py -- test_queued_tab_is_present_and_clickable
- * tests/e2e/test_evaluations.py -- test_in_progress_tab_is_present_and_clickable
- * tests/e2e/test_evaluations.py -- test_pending_review_tab_filters_correctly
- * tests/e2e/test_evaluations.py -- test_cancelled_tab_is_present_and_clickable
- * tests/e2e/test_evaluations.py -- test_status_tab_count_badges_are_numeric
- * tests/e2e/test_breadcrumb.py -- test_evaluations_breadcrumb_is_not_dashboard (xfail)
- * tests/e2e/test_breadcrumb.py -- test_models_breadcrumb_is_not_dashboard (xfail)
- * tests/e2e/test_breadcrumb.py -- test_breadcrumb_home_link_navigates_to_root
- * tests/e2e/test_breadcrumb.py -- test_breadcrumb_evaluation_workspace_link_works
- * tests/e2e/test_breadcrumb.py -- test_evaluation_detail_breadcrumb_has_org_name
- * tests/e2e/test_platform_issues.py -- test_meta_description_has_no_typo
- * tests/e2e/test_platform_issues.py -- test_meta_description_mentions_ai_evaluation
- * tests/e2e/test_platform_issues.py -- test_og_title_is_present
- * tests/e2e/test_platform_issues.py -- test_404_page_shows_not_found_message
- * tests/e2e/test_platform_issues.py -- test_404_page_has_home_navigation_link (xfail)
- * tests/e2e/test_platform_issues.py -- test_404_page_has_logo_or_branding (xfail)
- * tests/e2e/test_platform_issues.py -- test_hamburger_menu_is_visible_on_mobile
- * tests/e2e/test_platform_issues.py -- test_mobile_menu_dialog_has_accessible_title (xfail)
- * tests/e2e/test_platform_issues.py -- test_mobile_menu_has_navigation_links (xfail)
- * tests/e2e/test_platform_issues.py -- test_mobile_menu_can_be_closed_with_escape
- * tests/e2e/test_platform_issues.py -- test_unauthenticated_dashboard_redirects
- * tests/e2e/test_platform_issues.py -- test_unauthenticated_redirect_shows_login_prompt (xfail)
- * tests/e2e/test_platform_issues.py -- test_session_verification_completes_within_15s
- * tests/api/test_security.py -- test_html_page_has_security_headers
- * tests/api/test_security.py -- test_session_api_returns_empty_for_unauthenticated
- * tests/api/test_db_health.py -- test_graphql_responds_under_sequential_load
- * tests/api/test_db_health.py -- test_graphql_responds_under_5_concurrent_requests
- * tests/api/test_db_health.py -- test_graphql_error_response_does_not_contain_db_connection_error
- * tests/api/test_db_health.py -- test_models_endpoint_does_not_return_503
- * tests/api/test_db_health.py -- test_prompt_libraries_endpoint_does_not_return_503

5. Failure Details (First 20)

?

tests/accessibility/test_accessibility_auth.py::TestDashboardAccessibility::test_dashboard_has_exactly_one_h1[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/accessibility/test_accessibility_auth.py:78: in test_dashboard_has_exactly_one_h1
    assert count == 1, f"Expected 1 <h1> on dashboard, found {count}"
E   AssertionError: Expected 1 <h1> on dashboard, found 0
E   assert 0 == 1
```

? tests/e2e/test_evaluation_detail.py::TestEvaluationDetailSummary::test_summary_section_visible[chromium]

```
[gw2] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_evaluation_detail.py:54: in test_summary_section_visible
    assert ep.is_summary_section_visible()
E   assert False
E   + where False = is_summary_section_visible()
E   +           where is_summary_section_visible = <pages.evaluation_detail_page.EvaluationDetailPage object at 0x10b079910>.is_summary_section_visible
```

? tests/e2e/test_evaluation_detail.py::TestEvaluationDetailSummary::test_pass_rate_card_visible[chromium]

```
[gw2] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_evaluation_detail.py:58: in test_pass_rate_card_visible
    assert ep.is_pass_rate_visible()
E   assert False
E   + where False = is_pass_rate_visible()
E   +           where is_pass_rate_visible = <pages.evaluation_detail_page.EvaluationDetailPage object at 0x10f82c0d0>.is_pass_rate_visible
```

?

tests/accessibility/test_accessibility_auth.py::TestModelsPageAccessibility::test_model_cards_have_accessible_headings_or_labels[chromium]

```
arakhAI_test/.venv/bin/python3.11
tests/accessibility/test_accessibility_auth.py:236: in test_model_cards_have_accessible_headings_or_labels
    assert headings_in_cards.count() > 0, (
E   AssertionError: Model cards do not contain heading elements (h2/h3) or title elements. Screen readers need headings to navigate between cards.
E   assert 0 > 0
E   + where 0 = count()
E   +       where count = <Locator frame=<Frame name= url='https://dev.parakh.civicdataspace.in/dashboard/ai-maker/1/ai-models'> selector="[class*='card' i] h2, [class*='card' i] h3, [class*='card' i] [class*='title' i]">.count
```

? tests/api/test_security.py::TestCookieSecurity::test_session_cookies_have_httponly_flag[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/api/test_security.py:50: in test_session_cookies_have_httponly_flag
    assert not missing_http_only, (
E   AssertionError: Session cookies missing HttpOnly flag: ['KEYCLOAK_SESSION', 'KEYCLOAK_SESSION_LEGACY']. HttpOnly prevents JavaScript from accessing session cookies, mitigating XSS.
E   assert not ['KEYCLOAK_SESSION', 'KEYCLOAK_SESSION_LEGACY']
```

? tests/api/test_security.py::TestCookieSecurity::test_no_access_token_in_localstorage[chromium]

```
send(
  .venv/lib/python3.11/site-packages/playwright/_impl/_connection.py:69: in send
    return await self._connection.wrap_api_call(
  .venv/lib/python3.11/site-packages/playwright/_impl/_connection.py:559: in wrap_api_call
    raise rewrite_error(error, f"{parsed_st['apiName']}: {error}") from None
E   playwright._impl._errors.Error: Page.evaluate: SecurityError: Failed to read the 'localStorage' property from 'Window': Access is denied for this document.
E     at eval (<anonymous>)
E     at UtilityScript.evaluate (<anonymous>:302:30)
E     at UtilityScript.<anonymous> (<anonymous>:1:44)
```

? tests/api/test_security.py::TestCookieSecurity::test_no_password_in_localstorage[chromium]

```
send(
  .venv/lib/python3.11/site-packages/playwright/_impl/_connection.py:69: in send
    return await self._connection.wrap_api_call(
  .venv/lib/python3.11/site-packages/playwright/_impl/_connection.py:559: in wrap_api_call
    raise rewrite_error(error, f"{parsed_st['apiName']}: {error}") from None
E   playwright._impl._errors.Error: Page.evaluate: SecurityError: Failed to read the 'localStorage' property from 'Window': Access is denied for this document.
E     at eval (<anonymous>)
E     at UtilityScript.evaluate (<anonymous>:302:30)
E     at UtilityScript.<anonymous> (<anonymous>:1:44)
```

? tests/e2e/test_evaluation_detail.py::TestEvaluationDetailRiskCards::test_risk_section_visible[chromium]

```
[gw2] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_evaluation_detail.py:66: in test_risk_section_visible
    assert ep.is_risk_section_visible()
E   assert False
E   + where False = is_risk_section_visible()
E   +       where is_risk_section_visible = <pages.evaluation_detail_page.EvaluationDetailPage object at 0x10b75b110>.is_risk_section_visible
```

? tests/e2e/test_ai_maker_dashboard.py::TestModelCardsonHome::test_add_new_model_button_is_visible[chromium]

```
[gw1] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:191: in test_add_new_model_button_is_visible
    assert ai.is_add_new_model_visible(), "'Add A New Model' button must be visible"
E   AssertionError: 'Add A New Model' button must be visible
E   assert False
E   + where False = is_add_new_model_visible()
E   +       where is_add_new_model_visible = <pages.ai_maker_page.AIMakerPage object at 0x10901a8d0>.is_add_new_model_visible
```

? tests/e2e/test_ai_maker_dashboard.py::TestModelCardsonHome::test_model_cards_are_rendered[chromium]

```
[gw1] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:200: in test_model_cards_are_rendered
    assert count >= 1, f"Expected at least 1 model card on home, found {count}"
E   AssertionError: Expected at least 1 model card on home, found 0
E   assert 0 >= 1
```

? tests/e2e/test_ai_maker_dashboard.py::TestModelCardsonHome::test_text_generation_badge_is_present[chromium]

```
[gw1] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:209: in test_text_generation_badge_is_present
    assert found, "'Text Generation' badge must appear on model cards"
E   AssertionError: 'Text Generation' badge must appear on model cards
E   assert False
```

? tests/e2e/test_ai_maker_dashboard.py::TestOverviewStats::test_evaluation_runs_stat_is_visible[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:71: in test_evaluation_runs_stat_is_visible
    assert ai.is_stat_evaluation_runs_visible(), "Evaluation Runs stat must be visible"
E   AssertionError: Evaluation Runs stat must be visible
E   assert False
E   + where False = is_stat_evaluation_runs_visible()
E   +       where is_stat_evaluation_runs_visible = <pages.ai_maker_page.AIMakerPage object at 0x10d331b10>.is_stat_evaluation_runs_visible
```

? tests/e2e/test_ai_maker_dashboard.py::TestAIMakerButtons::test_add_organisation_button_is_visible[chromium]

```
i_maker_dashboard.py:221: in test_add_organisation_button_is_visible
    assert ai.is_visible(ai.ADD_ORGANISATION_BUTTON), (
E   AssertionError: 'Add Organisation' button must be visible on the AI Maker landing page
E   assert False
E   + where False = is_visible("button:has-text('Add Organisation')", a:has-text('Add Organisation'))
E   +       where is_visible = <pages.ai_maker_page.AIMakerPage object at 0x108ffb510>.is_visible
E   +       and "button:has-text('Add Organisation')", a:has-text('Add Organisation')" = <pages.ai_maker_page.AIMakerPage object at 0x108ffb510>.ADD_ORGANISATION_BUTTON
```

? tests/e2e/test_ai_maker_dashboard.py::TestOverviewStats::test_test_cases_stat_is_visible[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:76: in test_test_cases_stat_is_visible
    assert ai.is_stat_test_cases_visible(), "Test Cases stat must be visible"
E   AssertionError: Test Cases stat must be visible
E   assert False
E   + where False = is_stat_test_cases_visible()
E   +       where is_stat_test_cases_visible = <pages.ai_maker_page.AIMakerPage object at 0x10d332b50>.is_stat_test_cases_visible
```

? tests/e2e/test_evaluation_workspace.py::TestAIMakerRoleNavigation::test_org_selection_has_multiple_orgs[chromium]

ParakhAI Platform -- Full Test Report

```
[gw2] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_evaluation_workspace.py:128: in test_org_selection_has_multiple_orgs
    assert ws.is_visible(ws.AI_MAKER_CARD, timeout=5_000), (
E   AssertionError: AI Maker card not visible
E   assert False
E   +   where False = is_visible('text=AI Maker', timeout=5000)
E   +   where is_visible = <pages.workspace_page.WorkspacePage object at 0x10f818950>.is_visible
E   +   and   'text=AI Maker' = <pages.workspace_page.WorkspacePage object at 0x10f818950>.AI_MAKER_CARD
```

? tests/e2e/test_ai_maker_dashboard.py::TestOverviewStats::test_issues_flagged_stat_is_visible[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:86: in test_issues_flagged_stat_is_visible
    assert ai.is_stat_issues_flagged_visible(), "Issues Flagged stat must be visible"
E   AssertionError: Issues Flagged stat must be visible
E   assert False
E   +   where False = is_stat_issues_flagged_visible()
E   +   where   is_stat_issues_flagged_visible   = <pages.ai_maker_page.AIMakerPage object at 0x10d5cdad0>.is_stat_issues_flagged_visible
```

? tests/e2e/test_ai_maker_dashboard.py::TestOverviewStats::test_all_four_stat_cards_visible[chromium]

```
[gw0] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_ai_maker_dashboard.py:101: in test_all_four_stat_cards_visible
    assert not missing, f"Missing stat cards: {missing}"
E   AssertionError: Missing stat cards: ['Evaluation Runs', 'Test Cases', 'Issues Flagged']
E   assert not ['Evaluation Runs', 'Test Cases', 'Issues Flagged']
```

?

tests/e2e/test_evaluation_workspace.py::TestEvaluatorRoleNavigation::test_clicking_evaluator_navigates_to_evaluator_dashboard[chromium]

```
ct mode violation: locator("text=Evaluator") resolved to 2 elements:
E   1) <span class="Text-module_root__7LEnm Text-module_headingLg_K8tvY Text-module_semibold_EEhac text-gray-900 font-semibold text-lg sm:text-xl text-center">Evaluator</span> aka get_by_role("link", name="Evaluator Evaluator For")
E   2) <span class="Text-module_root__7LEnm Text-module_bodySm_1yBXn Text-module_regular__F5quw text-gray-600 text-center">For expert as evaluator</span> aka get_by_role("link", name="Evaluator Evaluator For")
E
E   Call log:
E   - waiting for locator("text=Evaluator") to be visible
```

? tests/e2e/test_auditor_evaluations.py::TestAuditorEvaluationsListPage::test_empty_state_or_rows_visible[chromium]

```
[gw1] darwin -- Python 3.11.15 /Users/home/CivicDataSpace/ParakhAI_test/.venv/bin/python3.11
tests/e2e/test_auditor_evaluations.py:52: in test_empty_state_or_rows_visible
    assert has_rows or has_empty, (
E   AssertionError: Neither evaluation rows nor empty-state message found on auditor evaluations page
E   assert (False or False)
```

? tests/e2e/test_evaluation_workspace_manual.py::TestManualWorkspaceLoad::test_can_navigate_to_wizard[chromium]

```
vigate_to_wizard
    assert ep.is_visible(ep.NEW_EVALUATION_BUTTON, timeout=3_000), (
E   AssertionError: New Evaluation button not visible -- platform may be unavailable or account lacks permission
E   assert False
E   +   where False = is_visible("button:has-text('New Evaluation')", a:has-text('New Evaluation')", timeout=3000)
E   +   where is_visible = <pages.evaluations_page.EvaluationsPage object at 0x10b338090>.is_visible
E   +   and   "button:has-text('New Evaluation')", a:has-text('New Evaluation')" = <pages.evaluations_page.EvaluationsPage object at 0x10b338090>.NEW_EVALUATION_BUTTON
```

6. Priority Recommendations

P0 -- Immediate

- * Deploy PgBouncer or increase PostgreSQL max_connections on dev (Finding #13)
- * Fix N+1 query on GetAudits/overview -- add DataLoader batching (Finding #5, #9)

- * Add server-side timeout on model list query; show error state in dialog (Finding #12)

P1 -- High Priority

- * Set Keycloak allowed-origins to specific domains, not '*' (Finding #15)
- * Add navigation links to mobile hamburger menu (Finding #19)
- * Fix breadcrumb label bug on sub-pages showing 'Dashboard' (Finding #7)
- * Create a branded custom 404 page (Finding #20)

P2 -- Medium Priority

- * Consider not exposing raw JWT in /api/auth/session client-side (Finding #14)
- * Add DialogTitle to mobile menu for screen reader accessibility (Finding #18)
- * Fix 'Completed on' column header for FAILED evaluations -- use 'Finished on' (Finding #8)
- * Correct meta description typo 'Paricipatory' ? 'Participatory' (Finding #1)
- * Show explicit login prompt when redirect occurs from protected URL (Finding #2)
- * Standardise URL locale prefix (/en/ vs no prefix) across all routes (Finding #3)

P3 -- Low Priority / Nice to Have

- * Make Evaluation Name read-only on COMPLETED evaluations (Finding #10)
- * Add evaluation name to breadcrumb on detail pages (Finding #11)
- * Hide pagination until data finishes loading (Finding #21)
- * Add Open Graph title/image tags for social sharing