


1. Executive Summary

A full Playwright MCP browser audit was conducted across all key pages of the ParakhAI platform on dev. **36 distinct issues** were identified across three audit sessions (core platform, Evaluator dashboard & invite flow, and the Add Model cross-platform flow). The desktop experience has a solid foundation but is undermined by two critical infrastructure failures (DB pool exhaustion and backend unavailability) that cause multiple pages to show indefinite "Loading..." states. These must be resolved before scaling concurrent usage or moving to production.

Infrastructure alert: Browser console logs expose connection to server at "127.0.0.1", port 5433 failed: FATAL: sorry, too many clients already during normal page loads. This reveals the internal database host, port, and a PostgreSQL FATAL error to any user with devtools open. Two categories of failure result: intermittent 503 errors on RSC prefetch routes, and indefinite "Loading..." states on AI Models, Prompt Libraries, Evaluators, and Evaluations pages.

Accessibility alert: The mobile hamburger menu dialog is broken in two ways — its aria-labelledby references a DOM element ID that does not exist (DialogTitle never rendered), and its content shows only a user avatar button with no navigation links. Mobile users cannot navigate the dashboard from the menu.



Responsive Coverage Matrix

Page / Feature	390px Mobile	768px Tablet	1440px Desktop
Homepage	✓ Pass — no overflow, fast load	✓ Pass — no overflow	✓ Pass
Navigation header	~ Hamburger only — no nav links in menu	~ Hamburger only (breakpoint >768px)	✓ Full nav visible
Dashboard (role select)	✓ Renders correctly	✓ Renders correctly	✓ Pass
AI Maker overview	✗ Hangs: "Loading overview..."	✗ Hangs: "Loading overview..."	✗ Hangs: "Loading overview..."
Evaluations list	✓ Table in scroll wrapper, no page overflow	✓ Pass	✓ Pass, 10 rows, 9 status tabs
Evaluation detail	✓ Loads, stats visible	✓ Pass	✓ Pass, Download Report present
AI Models	~ Loads after ~4s delay (intermittent)	~ Intermittent 503 on RSC prefetch	~ 503 on RSC prefetch; content loads eventually
Prompt Libraries	✗ "Loading prompt libraries..." (backend unavailable)	✗ 503 on RSC prefetch	✗ 503 on RSC prefetch
Evaluators	✗ "Loading evaluators..." (backend unavailable)	✗ "Loading evaluators..."	✗ "Loading evaluators..."
404 page	✗ Bare Next.js default — no nav or home link	✗ Bare Next.js default	✗ Bare Next.js default

2. Critical Issues

INFRA-001 · PostgreSQL connection pool exhaustion — DB credentials exposed in browser console Severity: CRITICAL Area: Infrastructure / Security Pages: AI Models, Evaluations, Prompt Libraries, Evaluators Browser console logs captured during this audit contain the exact PostgreSQL error: "connection to server at \"127.0.0.1\", port 5433 failed: FATAL: sorry, too many clients already" . This reveals the internal database host (127.0.0.1), the non-standard port (5433), and a fatal pool exhaustion event to any user who opens browser devtools. Concurrently, affected pages display "Loading..." indefinitely with no error message or retry option shown to the user. Confirmed in: AI Models fetch, Evaluations fetch, Prompt Libraries fetch, Auditors fetch. The pool is overwhelmed by concurrent GraphQL queries that each open a new connection — the N+1 query pattern on GetAudits (981+ rows) is a primary driver. Immediate: Deploy PgBouncer between the application and PostgreSQL to pool connections. Set max_connections appropriately and configure PgBouncer pool_mode=transaction. Medium-term: Implement DataLoader batching on the GetAudits resolver to eliminate N+1 queries. Security: Ensure GraphQL error messages are sanitised before being sent to the frontend — internal DB errors must not reach the browser console.
INFRA-002 · AI Maker overview "Loading overview..." hangs indefinitely (N+1 GetAudits) Severity: CRITICAL Area: Performance / Backend Page: AI Maker org home (/dashboard/ai-maker/1) After 5+ seconds the AI Maker org overview page still shows "Loading overview..." with no data rendered. This is the core landing page for the primary user role and is completely non-functional. The root cause is a GraphQL N+1 query: the overview resolver fetches 981+ audit rows individually rather than in a single batched query, exhausting the DB connection pool for other pages simultaneously. Confirmed to be unresolved since the previous MCP session (22 June 2026 session 1). Fix: Implement DataLoader batching on the GetAudits resolver. Replace per-row subqueries with a single JOIN query returning aggregated counts (total, by status). Add a server-side timeout with a cached fallback (e.g. 5-second timeout returns last known aggregate). Add a client-side timeout with a "Could not load overview — try refreshing" message instead of infinite spinner.

3. High Severity Issues

INFRA-003 · 503 Service Unavailable on RSC prefetch routes — still unresolved Severity: HIGH Area: Infrastructure Routes: /ai-models?_rsc=, /prompt-libraries?_rsc=, /en/dashboard/ai-maker/15?_rsc= RSC (React Server Component) prefetch requests consistently returnHTTP 503 for the AI Models and Prompt Libraries routes under org 1, and for org 15 (ASDMA). Confirmed in console: Failed to load resource: the server responded with a status of 503 on /dashboard/ai-maker/1/ai-models?_rsc=9y5yu and /dashboard/ai-maker/1/prompt-libraries?_rsc=9y5yu . These prefetch requests are triggered by Next.js on hover/navigation and are critical for fast page transitions. The 503s cause content to either show stale data or require full reload. Fix: Investigate the Next.js server and backend under RSC prefetch load. Add rate-limiting or concurrency controls on the RSC handler. Confirm that org 15 (ASDMA) has a backend data issue vs. a general routing problem — it 503s even at rest. Consider adding a circuit-breaker pattern on RSC routes.
AUTH-001 · Session token expires during active session — silent GraphQL failures Severity: HIGH Area: Authentication Component: SessionGuard Console warning captured: SessionGuard: Logging out - Session error: RefreshTokenExpired . This is followed by repeated SessionGuard: Backend validation error: TypeError: Failed to fetch errors that cause all GraphQL requests to fail silently — pages display "Loading..." with no explanation. The Keycloak refresh token has a short expiry and SessionGuard does not gracefully handle expiry during an active browsing session. Users are not prompted to re-authenticate; they are instead left with a broken dashboard that shows spinners indefinitely. Fix: Intercept RefreshTokenExpired in SessionGuard and show an in-page "Your session has expired — click to log in again" banner (not a silent logout). Implement proactive token refresh before expiry (e.g. at 80% of token lifetime using a setInterval check). Increase Keycloak refresh token TTL for the dev client.
SEC-001 · Internal database host and port exposed in browser console Severity: HIGH Area: Security / Information Disclosure The raw PostgreSQL error message including 127.0.0.1 (internal host) and port 5433 (non-standard dev port) reaches the browser console via the GraphQL error response chain. Additionally, the nginx server version (nginx/1.24.0 (Ubuntu)) is exposed in API response headers. Both reveal infrastructure details that aid an attacker in fingerprinting the stack. Fix GraphQL errors: Add a GraphQL error sanitiser/masker in Django middleware that replaces DB-level error messages with a generic "Service temporarily unavailable" before sending to the client. Fix nginx: Set server_tokens off; in nginx.conf to suppress version from response headers.
UX-001 · Mobile hamburger menu has no navigation links — dashboard inaccessible on mobile Severity: HIGH Area: UX / Mobile Viewport: 390px & 768px On mobile (390px) and tablet (768px) the "Open menu" hamburger button opens a Radix UI dialog whose entire content is a single SM (user avatar initials) button. There are no navigation links (Dashboard, AI Maker, Evaluations, Switch Roles, Logout). A mobile user who navigates away from the dashboard has no way to return without manually editing the URL. The desktop sidebar (Home, Models, Evaluations, Prompt Libraries, Evaluators) is hidden at all breakpoints below ~1024px with no mobile alternative. Fix: Populate the mobile menu dialog with the primary navigation links. At minimum: Dashboard, AI Maker (org select), current org sub-nav (if in org context), and Logout. Ensure the dialog also has a proper DialogTitle (see ACC-001).

4. Medium Severity Issues

UX-002 · All dashboard page titles identical: "Dashboard | ParakhAI"

Severity: MEDIUM | Area: SEO / UX | Pages: All dashboard routes

Every authenticated page (Evaluations, AI Models, Prompt Libraries, Evaluators, Eval Detail, AI Maker) has the same `<title>Dashboard | ParakhAI</title>`. Browser tabs are indistinguishable, back/forward navigation is ambiguous, and screen readers cannot orient users by page title. Also breaks bookmark management when multiple pages are bookmarked.

Fix: Set page-specific titles in `Next.js metadata exports`: "Evaluations — CivicDataLab | ParakhAI", "AI Models — CivicDataLab | ParakhAI", "Eval #1383 — 60% Pass Rate | ParakhAI", etc.

UX-003 · Breadcrumb items all show duplicated text ("HomeHome", "AI MakerAI Maker")

Severity: MEDIUM | Area: UX / Accessibility | Pages: All dashboard routes

Every breadcrumb item renders its text twice: "HomeHome", "Evaluation WorkspaceEvaluation Workspace", "AI MakerAI Maker", "CivicDataLabCivicDataLab". The duplication is visible in the accessibility tree and is announced doubled by screen readers. Root cause: the breadcrumb component likely renders two text nodes — one for visible text and one for an aria-label or screen-reader variant — without the second being hidden via `aria-hidden`.

Fix: Audit the `BreadcrumbItem` component. If there are two text spans, add `aria-hidden="true"` to the decorative copy. Alternatively, use a single text node and rely on CSS for styling without duplication.

UX-004 · Breadcrumb last item shows "Dashboard" not org/section name on sub-pages

Severity: MEDIUM | Area: UX | Pages: AI Models, Prompt Libraries, Evaluators, Eval Detail

On AI Models, Prompt Libraries, Evaluators, and Evaluation Detail pages the final breadcrumb segment shows "Dashboard" instead of the organisation name (CivicDataLab) or the section name (AI Models, Evaluators, etc.). The correct breadcrumb on the AI Models page should read: Home › Evaluation Workspace › AI Maker › CivicDataLab › AI Models.

Fix: Pass the correct `label` prop to `BreadcrumbItem` based on the route segment. For dynamic segments (org name, eval name) resolve via the page's server-side data fetch.

UX-005 · Missing H1 on 5 pages — AI Maker Select, Overview, Eval Detail, plus all sub-pages share generic title

Severity: MEDIUM | Area: Accessibility / SEO | WCAG 1.3.1

No H1 element was found on: AI Maker org select page, AI Maker org overview/home, or Evaluation Detail. Evaluations list has H1 "Evaluations"; AI Models has H1 "AI Models"; Evaluators has H1 "Evaluators" — these three are correct. Screen readers use H1 for page orientation. Without it, the page has no accessible landmark heading.

Fix: Add page-level H1s: "Select Organisation" (AI Maker select), "CivicDataLab Overview" (org home), and the evaluation name or "Evaluation #1383" (eval detail). H1 may be visually styled as a sub-heading if design requires.

UX-006 · Evaluation name is editable on COMPLETED evaluations

Severity: MEDIUM | Area: UX / Data Integrity | Page: Evaluation Detail (status: COMPLETED)

On evaluation 1383 (status: COMPLETED), the evaluation name input has `readOnly=false` and `disabled=false` — it is fully editable. Completed evaluations are final records; allowing name edits after completion creates an audit trail inconsistency. Users may accidentally rename a completed evaluation and lose context.

Fix: Set `readOnly={true}` or `disabled={true}` on the name input when `audit.status` is `COMPLETED`, `FAILED`, or `CANCELLED`. Apply the same guard to any other editable fields on the detail page.

UX-007 · Org card links use /en/ prefix but dashboard routes do not — inconsistent URL strategy

Severity: MEDIUM | Area: Routing / UX | Pages: AI Maker org select

The 10 organisation cards on the AI Maker select page all link to `/en/dashboard/ai-maker/{id}` (with locale prefix), but the sidebar navigation within the org uses `/dashboard/ai-maker/1` (without prefix). This inconsistency causes unnecessary 307 redirects when clicking org cards, breaks canonical URL strategy, and can cause hydration mismatches.

Fix: Standardise all internal links to either always include or never include the `/en/` prefix. Use `Next.js useRouter` with the `locale` option to generate locale-aware paths consistently.

UX-008 · Pagination controls render before data loads ("Page 01 of 01" during fetch)

Severity: MEDIUM | Area: UX | Pages: AI Models, Prompt Libraries

On AI Models and Prompt Libraries, the pagination footer ("Results per page / Page 01 of 01 / First · Prev · Next · Last") renders immediately while the data is still loading. This is misleading — "Page 01 of 01" implies no results exist when the fetch is in flight. After load, AI Models correctly shows "Page 01 of 02".

Fix: Conditionally render the pagination component only after the data query resolves (i.e. when `isLoading === false`). Show a spinner or skeleton in place of pagination during load.

UX-009 · Table sort column headers missing aria-sort attribute

Severity: MEDIUM | Area: Accessibility | WCAG 1.3.1 | Page: Evaluations table

The 7 sortable column headers in the Evaluations table (Evaluation Name, Model, Evaluation Type, Status, Evaluation Mode, Tests, Completed on) display "Sorted in descending order" as visible text but have no `aria-sort` attribute. Screen readers cannot programmatically determine the current sort direction or whether a column is sortable. The sort buttons also lack descriptive `aria-label` values identifying which column they sort.

Fix: Add `aria-sort="descending"` (or "ascending") to the active sort column's `<th>`. Add `aria-sort="none"` to unsorted columns. Add `aria-label="Sort by Evaluation Name"` to each sort button.

UX-010 · New Evaluation dialog aria-describedby references a non-existent DOM element

Severity: MEDIUM | Area: Accessibility | Page: Evaluations (New Evaluation dialog)

The New Evaluation dialog sets `aria-describedby="radix:r6:"` but no element with `id="radix:r6:"` exists in the DOM (the `DialogDescription` component was never rendered). This is a broken ARIA reference — the attribute points to nothing, which is equivalent to having no description. Screen readers may announce nothing or

skip the description entirely. Console warning confirms: `Warning: Missing Description or aria-describedby={undefined} for {DialogContent}` .

Fix: Add a `<DialogDescription>` element to the New Evaluation dialog (e.g. "Select a model and configure your evaluation run."). Even if visually hidden, it must exist in the DOM for the `aria-describedby` reference to be valid.

UX-011 · Form selects in New Evaluation dialog have no accessible label

Severity: MEDIUM | Area: Accessibility | WCAG 1.3.1 | Page: New Evaluation dialog

The "Select an AI Model" (`name="modelSelect"`) and version (`name="versionSelect"`) dropdowns have no `aria-label` , no `aria-labelledby` , and their associated `<label>` elements are present in the DOM but not linked via `for` / `id` pairs. Screen readers announce these as unlabelled comboboxes.

Fix: Ensure each `<label>` has a `for` matching the input's `id` , or add `aria-labelledby` to each input pointing to the label element. Verify with an axe audit after fixing.

ACC-001 · Mobile menu DialogTitle not rendered — broken aria-labelledby reference

Severity: MEDIUM | Area: Accessibility | WCAG 4.1.2 | Component: Mobile menu dialog

The mobile hamburger menu dialog has `aria-labelledby="radix:r1:"` but no element with that ID exists in the DOM. The Radix UI `DialogTitle` was never rendered, so the reference is broken. Console error: `DialogContent requires a DialogTitle for the component to be accessible for screen reader users` . Screen readers cannot determine the dialog's purpose.

Fix: Add `<DialogTitle>` to the mobile menu dialog. If no visual title is desired, wrap it in Radix's `<VisuallyHidden>` : `<VisuallyHidden><DialogTitle>Navigation Menu</DialogTitle></VisuallyHidden>` .

5. Low Severity Issues

SEO-001 · Meta description contains typo "Paricipatory" (missing 't') Severity: LOW Area: SEO Page: All pages (global meta) The site-wide meta description is "Paricipatory AI Evaluation" — "Participatory" is misspelled. This appears in Google search result snippets and is the public-facing description of the platform. Fix: Correct to "Participatory" in the root <code>layout.tsx</code> metadata export. While there, expand the description to include the platform's purpose (model registry, audit, evaluation) for better search visibility.
SEO-002 · No Open Graph meta tags — social sharing broken Severity: LOW Area: SEO / Social Page: All pages <code>og:title</code> , <code>og:description</code> , and <code>og:image</code> are all absent. When the platform URL is shared on LinkedIn, Slack, WhatsApp, or Twitter/X, no preview card is generated — only the bare URL is shown. Fix: Add Open Graph and Twitter Card meta tags in <code>layout.tsx</code> : <code>openGraph: { title, description, images: [{ url: '/og-image.png' }] }</code> . Create a branded <code>og-image.png</code> (1200x630px).
SEO-003 · No canonical link tag on any page Severity: LOW Area: SEO Page: All pages No <code><link rel="canonical"></code> tag is present. Combined with the <code>/en/</code> prefix inconsistency (INFRA issue), the same page may be indexed at two different URLs (<code>/dashboard/ai-maker/1</code> and <code>/en/dashboard/ai-maker/1</code>) causing duplicate content penalties. Fix: Add <code>alternates: { canonical: '...' }</code> to each page's metadata export. Resolve the locale prefix inconsistency first (UX-007) to determine the canonical URL pattern.
UX-012 · 404 page is bare Next.js default — no branding or recovery path Severity: LOW Area: UX Route: Any invalid URL Any invalid URL (e.g. <code>/this-page-does-not-exist</code> , <code>/login</code>) renders the default Next.js 404: white page, "404 / This page could not be found." , system font, inline styles only. No header, no logo, no navigation, no "Go to Dashboard" or "Go Home" link. Users hitting broken links are completely stranded. Fix: Create <code>app/not-found.tsx</code> (or <code>app/[locale]/not-found.tsx</code>) rendering the full site shell with: ParakhAI logo, nav bar, "Page not found" message, and a prominent "Go to Dashboard" CTA button.
PERF-001 · Unused CSS preload: <code>c7b7feaf3a90da3e.css</code> preloaded but not consumed Severity: LOW Area: Performance Page: AI Models Console warning on the AI Models page: The resource <code>.../c7b7feaf3a90da3e.css</code> was preloaded using link preload but not used within a few seconds from the window's load event. Unused preloads consume bandwidth and a browser connection slot without benefit, competing with resources that are actually needed. Fix: Audit the Next.js build output for this chunk. If it belongs to a code-split route, ensure the preload is only injected on pages that use it. Review <code>next.config.js</code> for overly aggressive preload directives.
SEC-002 · No Content-Security-Policy header on the frontend Severity: LOW Area: Security Scope: All pages A <code>Content-Security-Policy</code> header is absent from all frontend responses. CSP is the primary defence against XSS attacks that bypass input sanitisation. All other major security headers (HSTS, <code>x-frame-options</code>, <code>x-content-type-options</code>, <code>referrer-policy</code>, <code>permissions-policy</code>) are correctly set. CSP is the only significant gap. Fix: Start with <code>Content-Security-Policy-Report-Only</code> mode to identify violations without breaking the site. Then enforce progressively. Configure via <code>Next.js headers()</code> in <code>next.config.js</code> or <code>nginx</code>. Baseline: <code>default-src 'self'; script-src 'self' 'nonce-{...}'</code> .
SEC-003 · API server version exposed in response headers (nginx/1.24.0 Ubuntu) Severity: LOW Area: Security / Information Disclosure Scope: API server The API server (<code>dev.api.parakh.civicdataspace.in</code>) returns <code>server: nginx/1.24.0 (Ubuntu)</code> in response headers. This reveals the exact nginx version and OS distribution, enabling targeted vulnerability scanning. The frontend server only exposes <code>server: nginx</code> (version suppressed — correct). Fix: Add <code>server_tokens off;</code> to the API server's <code>nginx.conf</code> to suppress the version number. Apply to all server blocks.

6. Security Test Results

Overall posture: Good baseline, two gaps. All standard security headers are present on the frontend. XSS reflection is safe. The two issues are: (1) no Content-Security-Policy, and (2) internal DB error details reaching the browser console.

Test	Method / Payload	Result	Status
Reflected XSS — query param	?search=%3Cscript%3Ealert(1)%3C/script%3E	URL-encoded in query string. Not injected into DOM. No alert executed. Page renders normally.	SAFE
DOM XSS — body reflection check	JS: check body.innerText and body.innerHTML for "alert(1)"	Not reflected in visible text or HTML structure. Next.js correctly escapes query parameter values.	SAFE
Unauthenticated API access	curl /graphql/ without token	Returns HTTP 403. Correct — GraphQL endpoint requires authentication.	SAFE
HSTS header	curl -I frontend	strict-transport-security: max-age=31536000; includeSubDomains; preload — correct and strong.	PASS
Clickjacking protection	curl -I frontend	x-frame-options: SAMEORIGIN — present on both frontend and API.	PASS
MIME sniffing	curl -I frontend	x-content-type-options: nosniff — present.	PASS
Referrer leakage	curl -I frontend	referrer-policy: strict-origin-when-cross-origin — correct.	PASS
Permissions policy	curl -I frontend	permissions-policy: geolocation=(), microphone=(), camera=() — present.	PASS
Content-Security-Policy	curl -I frontend	Header absent on frontend. No CSP protection against script injection.	MISSING
DB error information exposure	Browser console during AI Models load	Raw PostgreSQL error: 127.0.0.1:5433 FATAL: sorry, too many clients visible in browser devtools.	FAIL
Server version exposure (API)	curl -I API server	server: nginx/1.24.0 (Ubuntu) — version exposed. Frontend correctly shows only "nginx".	MINOR

7. Performance Metrics (Navigation Timing API)

All time-to-first-byte and DOM load metrics are excellent (<120ms). The performance problem is not network or render speed — it is backend query latency (N+1, DB pool exhaustion) causing application-level loading spinners that persist for seconds or indefinitely.

Page	URL	TTFB (ms)	DOMContentLoaded (ms)	Load Complete (ms)	Notes
Homepage	/	41	103	122	Clean. No console errors.
Dashboard (role select)	/dashboard	43	79	80	Fast. Correct H1 "Select Your Role".
AI Maker org select	/dashboard/ai-maker	43	113	117	10 orgs listed. Breadcrumb duplicated.
AI Maker overview	/dashboard/ai-maker/1	39	109	113	Overview content never loads (N+1 hang).
Evaluations list	/dashboard/ai-maker/1/evaluations	45	76	82	10 rows loaded after ~4s. 9 status tabs.
Evaluation detail	/dashboard/ai-maker/1/evaluations/1383	44	110	110	Loads after ~5s. 60% pass, stats visible.
AI Models	/dashboard/ai-maker/1/ai-models	43	76	79	503 on RSC prefetch. Content loads ~4s.
Prompt Libraries	/dashboard/ai-maker/1/prompt-libraries	44	78	81	503 on RSC. Content did not load (backend unavailable).
Evaluators	/dashboard/ai-maker/1/auditors	45	74	103	Content did not load (backend unavailable).
Mobile homepage (390px)	/	44	72	83	No overflow. H1 correct. Hamburger present.

Key Observation

Network TTFB is consistently **39–45ms** — the server is responding fast. DOMContentLoaded is **72–113ms** — the initial HTML shell renders quickly. The UX performance problem is entirely on the **GraphQL data fetch layer**: pages render loading skeletons immediately, then wait 4–10+ seconds for backend queries, or hang indefinitely when the DB pool is exhausted. Fixing the DB connection pool and N+1 queries will have a larger user impact than any frontend optimisation.

8. What Works Well — Verified Functionality

✓ Core Data Flows (Desktop)

- Homepage renders correctly with H1 and CTA buttons
- Login button redirects to Keycloak SSO correctly
- Dashboard role selector (AI Maker) works
- AI Maker org select lists all 10 orgs with links
- Evaluations table loads 10 rows with correct columns
- 9 status filter tabs (All, Draft, Queued...) all present
- Evaluation detail: stats, pass rate, Download Report button
- AI Models: 9 cards, Page 1 of 2, model names and descriptions
- New Evaluation dialog: model dropdown loads immediately (no hang)
- New Evaluation: Back/Next wizard navigation works
- Sidebar nav: Home, Models, Evaluations, Prompt Libraries, Evaluators
- Table horizontal scroll wrapper at mobile (1457px in 375px container)

✓ Security & Infrastructure (Positive)

- HSTS: max-age=31536000, includeSubDomains, preload
- x-frame-options: SAMEORIGIN on frontend and API
- x-content-type-options: nosniff on both servers
- x-xss-protection: 1; mode=block
- referrer-policy: strict-origin-when-cross-origin
- permissions-policy: geo, mic, camera all blocked
- GraphQL API requires auth — 403 for unauthenticated
- XSS query param injection: safe, not reflected in DOM
- Mobile homepage: no horizontal overflow at 390px
- Tablet homepage: no horizontal overflow at 768px
- Hamburger menu: accessible label "Open menu"
- Homepage images: all have alt attributes

✓ AI Models Registered in System (sample from page 1 of 2)

xAI: Grok 4.3 — 613 test cases, 29 evaluations
Google: Gemma 4 26B A4B
SarvamAI: Sarvam 30B
MoonshotAI: Kimi K2.5
xAI: Grok 4.1 Fast (2M context)
Anthropic: Claude 3.5 Haiku
MistralAI: Mistral Large 3 2512
SarvamAI: Sarvam-M
OpenAI: GPT-5 Mini

✓ Evaluations Status Distribution (live data from Evaluations tab)

All: 100 total evaluations
Draft: 40
Queued: 0
Running: 1
In Progress: 10
Pending Review: 17
Completed: 25
Failed: 6
Cancelled: 1

9. Evaluator Dashboard & Invite Auditor Flow

The Evaluator role dashboard (/dashboard/auditor) was tested separately from the AI Maker dashboard. Three sub-pages exist: Overview (home), Assigned Models, Evaluations History. The full invite-evaluator flow (Add Evaluator → email search → confirm) was executed end-to-end. **The invite flow works correctly** but has a silent success bug. Four new issues were found.

Evaluator Dashboard — Pages & Status

Page	URL	H1	Loads	Notes
Evaluator Home (Overview)	/dashboard/auditor	✓ (none, but stats visible)	✓ ~4s	Stats: 794 evals completed, 10345 test cases, 6203 issues flagged. No pending invitations.
Assigned Models	/dashboard/auditor/assignments	✓ "Assigned Models"	✓ Fast	3 status tabs: All / Accepted(0) / Declined(0). Empty state is well-handled with helpful copy.
Evaluations History	/dashboard/auditor/evaluations	✓ "Evaluations History"	✓ Fast	Same 9-tab pattern as AI Maker. Empty state shows "View Assignments" CTA — good UX.

What Works Well — Evaluator Dashboard

- Breadcrumb correctly uses aria-current="page" on the last item (not the case on AI Maker side)
- Sidebar nav labels are descriptive: Home, Assigned Models, Evaluations History
- Overview stats (completions, test cases, issues) load correctly after ~4s
- Empty state handling on Assigned Models and Evaluations History is clear and actionable
- Evaluations History "View Assignments" CTA links correctly to /dashboard/auditor/assignments

Invite Evaluator Flow (Add Evaluator from AI Maker → Evaluators page)

Step	Action	Result	Status
1	Click "Add Evaluator" button	Dialog opens immediately with email input	PASS
2	Type email: saqibmanan6@gmail.com	Email accepted in input field	PASS
3	Click Search	User found: "saqib manan · saqibmanan6@gmail.com" shown as result card	PASS
4	Search button after result	Search button auto-disables after result found (correct)	PASS
5	Click "Add Evaluator" (confirm)	Button changes to "Adding..." (disabled) — in-progress state shown	PASS
6	Mutation completes	Dialog closes. New evaluator appears in list without page reload.	PASS
7	Success feedback	No toast, banner, or confirmation message shown. Silent success.	MISSING
8	Evaluator list updated	"saqib manan" now appears in list with Remove button — optimistic update works	PASS

New Issues Found — Evaluator Dashboard & Invite Flow

INFRA-004 · URL construction bug: "civicdataspace.innull" — null env var concatenated to API base URL

Severity: HIGH | Area: Infrastructure / Code Bug | Page: Evaluators (/dashboard/ai-maker/1/auditors)

6 network requests were recorded failing with ERR_NAME_NOT_RESOLVED targeting https://dev.api.civicdataspace.innull/. The string "null" is being appended directly to the API base URL, indicating a JavaScript string concatenation where a variable or environment value is null or undefined :

"https://dev.api.civicdataspace.in" + null → "https://dev.api.civicdataspace.innull" . These requests are triggered specifically on the Evaluators page and are not present on other pages. Likely a path suffix or query parameter variable resolving to null before being appended to the URL.

Fix: Search the frontend codebase for Config.url , process.env.NEXT_PUBLIC_* , or API client URL construction in the auditors page component. Add a guard: if (!pathVar) throw new Error('path is null') before concatenating. Also validate all env vars at app startup.

UX-013 · No success feedback after "Add Evaluator" — silent mutation

Severity: MEDIUM | Area: UX | Page: Evaluators (Add Evaluator dialog)

When an evaluator is successfully added, the dialog closes without showing any confirmation toast, banner, or status message. The only indication of success is that the evaluator appears in the list (which a user may not notice if the list is long). No equivalent error feedback was tested — an "already an evaluator" or "user not found on CivicDataSpace" scenario is also likely silent. Consistent with the broader pattern of silent state mutations across the platform.

Fix: Show a toast notification on success ("Evaluator added successfully") and on error ("User is already an evaluator" / "No user found with that email"). The platform already has a Notifications region (region "Notifications alt+T" in the ARIA tree) — wire the mutation result into it.

UX-014 · Add Evaluator dialog heading renders twice in DOM

Severity: MEDIUM | Area: UX / Accessibility | Page: Add Evaluator dialog

The dialog snapshot shows two heading "Add Evaluator" [level=2] nodes inside the dialog — at separate ref positions (e344 and e347). One is the Radix UI DialogTitle (correct); the other is a visual heading rendered inside the content area (duplicate). Screen readers announce "Add Evaluator" twice when focus enters the dialog. This is the same root cause as the breadcrumb duplication — a shared component renders both a hidden accessible title and a visible heading without deduplication.

Fix: Use the Radix <DialogTitle> as the single source of truth for the heading. Remove the duplicate <h2> from the dialog body content, or mark it aria-hidden="true" .

DATA-001 · Evaluator list data quality: typo "Uday Raj12 Sahai" and duplicate entry

Severity: LOW | Area: Data Quality | Page: Evaluators list

The evaluators list for org CivicDataLab contains two data quality issues: (1) "Uday Raj12 Sahai" has a number "12" embedded in the name — likely a testing artefact or data entry error. (2) "Uday Raj Sahai" appears twice with different designations (once with no title, once as "Frontend Developer") — suggesting duplicate user registration or re-invitation of the same person.

Fix: (1) Correct the name via the Manage Users interface or directly in the database. (2) Deduplicate the evaluator entry — investigate whether two separate user accounts exist for the same person and merge them if so.

10. Add Model Flow (ParakhAI → CivicDataSpace)

The "Add A New Model" flow on ParakhAI redirects to a 3-step wizard on CivicDataSpace. This section documents the full cross-platform journey and the new issues found.

Flow Summary

Step 1 — Trigger: ParakhAI AI Models page → "Add A New Model" button → redirect confirmation dialog opens.

Step 2 — Redirect: "Yes, continue" opens a *new tab* to CivicDataSpace (/dashboard/organization/civicaldatallab/aimodels?tab=registered). User is already authenticated via shared Keycloak session.

Step 3 — CivicDataSpace wizard (3 tabs): Metadata → Version → Publish. The Publish tab is the only validation gate; tabs 1→2 can be traversed freely with empty required fields.

Step 4 — Version creation: "Create First Version" dialog creates a version record (v1.0, Development lifecycle).

Step 5 — Access Method: "Add Access Method" dialog configures API provider (OpenAI, Ollama, Together AI, Replicate, Custom) with endpoint URL, API key (password field), auth type, request/response body templates, and timeout.

HIGH — Step 1→2 Bypasses All Required Field Validation

VALID-001 · Step 1→2 wizard navigation skips all required field validation

Clicking "Next" from the Metadata step advances to the Version step without validating any required field (Sectors*, Tags*, Languages*, Geographies*, Intended Use*, Target Users*, Model Website*, Maximum Tokens*). Only the Publish step (step 3) enforces completeness — the Publish button is disabled and a list of missing fields is shown. A user can create a Version and configure Access Methods on a model with zero metadata, and only discover the block at the very last step. Risk: orphaned version/access-method records attached to an incomplete model draft.

Fix: Add per-step validation before advancing: step 1 → 2 requires all Metadata required fields to be non-empty; step 2 → 3 requires at least one Access Method to be configured.

HIGH — CivicDataSpace JS SyntaxError on Every Editor Page Load

CDS-001 · SyntaxError: Failed to execute 'appendChild' on 'Node': missing) after argument list

A SyntaxError fires on every page load of the CivicDataSpace AI model editor (chunk549-3cc5a4b2f9dfcb5c.js). The error originates from a p() call inside the compiled Next.js bundle and propagates through the React fiber renderer. The error fires at least twice per page load. While the page partially renders, the error may cause incomplete renders of specific UI sub-trees on slower connections or lower-powered devices.

Fix: Locate the source of the dynamic appendChild call in the unminified source — likely a template literal or dangerouslySetInnerHTML block with an unescaped parenthesis. Rebuild and redeploy the affected chunk.

MEDIUM — Ghost Model Created Before Any Form Input

CDS-002 · Clicking "Add New AI Model" immediately creates a DB record (no form shown first)

On the CivicDataSpace AI models page, clicking "Add New AI Model" immediately creates a model record in the database (e.g. ID 168, named "New AI Model") and navigates to the edit wizard. If the user closes the editor without filling any data, an orphaned draft model with no name, no metadata, and no version remains in the database indefinitely. This pollutes the Draft list and may count against org quotas.

Fix: Either (a) require a model name before creating the DB record (a simple "Name your model" modal), or (b) implement auto-deletion of models that are closed within 60 seconds of creation with zero metadata filled.

MEDIUM — Autosave Indicator Misleads: "All Changes Saved" With Empty Required Fields

CDS-003 · "All Changes Saved" indicator shows even when all required metadata fields are empty

The editor header always shows "All Changes Saved" regardless of whether required fields have been populated. The autosave tracks the model record's existence, not field completeness. This creates a false sense of security — a user who fills nothing sees "All Changes Saved" and may believe their work is complete. Separately confirmed: autosave fires GraphQL mutations on every change event, but the indicator does not distinguish "draft record saved" from "complete, valid model saved".

Fix: Change the autosave indicator to show "Draft saved — metadata incomplete" when required fields are empty. Only show "All Changes Saved" when the model could pass the Publish validation.

MEDIUM — All "View" Buttons on Model Cards Are Disabled

UX-015 · Model card "View" buttons rendered as [disabled] — cannot view individual model details

On the ParakhAI AI Models list page (/dashboard/ai-maker/1/ai-models), every model card has a "View" button that is in a disabled state. There is no error message or tooltip explaining why. This means there is no way to navigate to individual model detail pages from the list. The card area itself is clickable ([cursor=pointer]) but clicking it does nothing observable — navigation does not occur.

Fix: Enable the View button when the model has at least one access method configured. If the feature is intentionally gated (e.g. role-based), show a tooltip explaining why ("Request access to view this model"). Remove the disabled state for models the user owns.

LOW — ?tab=registered URL Param from ParakhAI Silently Ignored

CDS-004 · Redirect URL includes ?tab=registered but no such tab exists

The ParakhAI redirect sends users to /dashboard/organization/civicaldatallab/aimodels?tab=registered. The CivicDataSpace AI models page only has "Draft" and "Published" tabs. The registered param is silently discarded and the page defaults to the Draft tab. If "registered" was intended to map to "Published", the param value is wrong; if it mapped to a planned future tab, that tab has not been built yet.

Fix: Align the redirect URL to use ?tab=draft or ?tab=published matching the actual tab values, or remove the param entirely.

LOW — API Key Field Has autocomplete="on"

SEC-004 · "Add Access Method" API key password field missing autocomplete="off"

The API key input in the "Add New Access Method" dialog is atype="password" field with autocomplete="on". Browser password managers may cache the value keyed to the current URL, and the key could autofill on unintended pages or be stored in browser sync. API keys should opt out of autocomplete to prevent accidental exposure.

Fix: Set autocomplete="off" (or autocomplete="new-password" to prevent browsers from overriding the attribute) on the API key input.

LOW — Quill bullet Format Not Registered

CDS-005 · Quill rich text editor: Cannot register "bullet" specified in "formats" config

The description rich text editor on the Metadata step logs: quill Cannot register "bullet" specified in "formats" config. Are you sure it was registered twice on load. The bullet list type is declared in the formats config but has not been registered with Quill's module system. This will cause the bullet-list toolbar button to silently fail when clicked.

Fix: Either register the bullet Quill blot before instantiating the editor, or remove it from the formats config if bullet lists are not required.

11. Master Issue Tracker — All 36 Issues

#	Title	Page / Area	Severity	Category
INFRA-001	DB connection pool exhaustion — FATAL error exposed in console	AI Models, Evaluations, APIs	CRITICAL	Infrastructure
INFRA-002	AI Maker overview "Loading overview..." hangs indefinitely (N+1)	AI Maker org home	CRITICAL	Performance
INFRA-003	503 on RSC prefetch — /ai-models and /prompt-libraries routes	AI Models, Prompt Libraries	HIGH	Infrastructure
AUTH-001	Session token expiry causes silent GraphQL failures	All authenticated pages	HIGH	Authentication
SEC-001	Internal DB host (127.0.0.1:5433) exposed in browser console	AI Models, API errors	HIGH	Security
UX-001	Mobile hamburger menu has no navigation links	All pages (mobile/tablet)	HIGH	UX / Mobile
UX-002	All dashboard page titles identical: "Dashboard ParakhAI"	All dashboard routes	MEDIUM	SEO / UX
UX-003	Breadcrumb text duplicated ("HomeHome", "AI MakerAI Maker")	All dashboard routes	MEDIUM	UX / Accessibility
UX-004	Breadcrumb last item shows "Dashboard" not org/section name	AI Models, Prompt Libs, Evaluators, Eval Detail	MEDIUM	UX
UX-005	Missing H1 on AI Maker Select, Overview, Eval Detail	3 pages	MEDIUM	Accessibility / SEO
UX-006	Eval name editable on COMPLETED evaluations	Evaluation Detail	MEDIUM	UX / Data Integrity
UX-007	Org links use /en/ prefix; sidebar uses non-prefixed URLs	AI Maker org select	MEDIUM	Routing
UX-008	Pagination renders before data loads ("Page 01 of 01" during fetch)	AI Models, Prompt Libraries	MEDIUM	UX
UX-009	Table sort headers missing aria-sort attribute	Evaluations table	MEDIUM	Accessibility
UX-010	New Eval dialog aria-describedby references non-existent element	New Evaluation dialog	MEDIUM	Accessibility
UX-011	Form selects in New Eval dialog have no accessible label	New Evaluation dialog	MEDIUM	Accessibility
ACC-001	Mobile menu DialogTitle not rendered — broken aria-labelledby	Mobile menu (all pages)	MEDIUM	Accessibility
SEO-001	Meta description typo: "Paricipatory" → "Participatory"	All pages	LOW	SEO
SEO-002	No Open Graph meta tags (og:title, og:description, og:image)	All pages	LOW	SEO / Social
SEO-003	No canonical link tag on any page	All pages	LOW	SEO
UX-012	404 page is bare Next.js default — no nav or home link	All invalid routes	LOW	UX
PERF-001	Unused CSS preload: c7b7feaf3a90da3e.css	AI Models page	LOW	Performance
SEC-002	No Content-Security-Policy header	All pages (frontend)	LOW	Security
SEC-003	API server exposes nginx version (nginx/1.24.0 Ubuntu)	API server headers	LOW	Security
INFRA-004	"civicdataspace.innull" — null env var concatenated to API URL	Evaluators page	HIGH	Infrastructure
UX-013	No success feedback after Add Evaluator — silent mutation	Add Evaluator dialog	MEDIUM	UX
UX-014	Add Evaluator dialog heading renders twice in DOM	Add Evaluator dialog	MEDIUM	UX / Accessibility
DATA-001	Evaluator name typo "Uday Raj12 Sahai" + duplicate entry	Evaluators list	LOW	Data Quality
VALID-001	Step 1 → 2 wizard navigation skips all required field validation	CivicDataSpace model editor	HIGH	Validation
CDS-001	JS SyntaxError appendChild missing) on every editor page load	CivicDataSpace model editor	HIGH	Infrastructure
CDS-002	Ghost model DB record created before any form input on "Add New AI Model"	CivicDataSpace AI Models	MEDIUM	Data Integrity

#	Title	Page / Area	Severity	Category
CDS-003	"All Changes Saved" shown with empty required fields — misleading autosave	CivicDataSpace model editor	MEDIUM	UX
UX-015	All "View" buttons on model cards disabled — can't view model details	AI Models list (ParakhAI)	MEDIUM	UX
CDS-004	?tab=registered URL param from ParakhAI redirect silently ignored	CivicDataSpace AI Models	LOW	Routing
SEC-004	API key field has autocomplete="on" in Add Access Method dialog	CivicDataSpace model editor	LOW	Security
CDS-005	Quill bullet format not registered — bullet list button silently fails	CivicDataSpace model editor	LOW	Editor / UX

Recommended Sprint Plan

Sprint 1 — Infrastructure (block concurrent use): INFRA-001 (PgBouncer + GraphQL error masking), INFRA-002 (DataLoader batching on GetAudits), AUTH-001 (session refresh + expiry UX), INFRA-003 (503 RSC investigation), CDS-001 (JS SyntaxError in model editor)

Sprint 2 — UX Blockers (before wider rollout): UX-001 (mobile nav links), UX-012 (custom 404), UX-003/004 (breadcrumb fixes), UX-002 (page titles), SEC-001 (mask DB errors in GraphQL), VALID-001 (step validation), UX-015 (View buttons disabled), CDS-002 (ghost model), INFRA-004 (civicdataspace.innull)

Sprint 3 — Accessibility & data integrity: ACC-001, UX-009, UX-010, UX-011, UX-005 (H1s), UX-006 (read-only on COMPLETED), CDS-003 (autosave indicator), UX-013/014 (Radix dialog heading duplication)

Sprint 4 — SEO, security & polish: SEO-001 (typo), SEO-002 (OG tags), SEO-003 (canonical), UX-007 (URL strategy), UX-008 (pagination skeleton), PERF-001, SEC-002 (CSP), SEC-003 (nginx version), SEC-004 (API key autocomplete), CDS-004 (tab param mismatch), CDS-005 (Quill bullet)